

09. – 12.12.2019
Frankfurt am Main



Bernhard Hopfenmüller

Apache Kafka -

- a system optimized for writing

#ittage

Ever named a software project?

Voldemort - An evil wizard



Source: Landmark Media/ImageCollect
Project: <https://github.com/voldemort/voldemort>

Jay Kreps:

"I don't know whether it is nerdier to be reading Harry Potter or to be wondering what kind of consistency protocol Voldemort uses when keeping all his pieces up-to-date, but regardless, the name stuck."

Voldemort - A distributed key value store



Source: Landmark Media/ImageCollect
Project: <https://github.com/voldemort/voldemort>

Jay Kreps:

"I don't know whether it is nerdier to be reading Harry Potter or to be wondering what kind of consistency protocol Voldemort uses when keeping all his pieces up-to-date, but regardless, the name stuck."

Azkaban - The wizard prison



Source: <https://www.bustle.com>
Project: <https://azkaban.github.io/>

Richard Park:

[...] clearly a phase where we started for some reason naming things after Harry Potter stuff [...]

Azkaban - A Distributed Workflow Manager



Source: <https://www.bustle.com>
Project: <https://azkaban.github.io/>

Richard Park:

[...] clearly a phase where we started for some reason naming things after Harry Potter stuff [...]



Jay Kreps:

I thought that since Kafka was a system optimized for writing using a writer's name would make sense. I had taken a lot of lit classes in college and liked Franz Kafka. Plus the name sounded cool for an OS project.

Kafka - A system optimized for writing



Jay Kreps:

I thought that since Kafka was a system optimized for writing using a writer's name would make sense. I had taken a lot of lit classes in college and liked Franz Kafka. Plus the name sounded cool for an OS project.

Kafka - A (distributed) system ...



Jay Kreps:

I thought that since Kafka was a system optimized for writing using a writer's name would make sense. I had taken a lot of lit classes in college and liked Franz Kafka. Plus the name sounded cool for an OS project.

curl -X GET localhost://whoami



```
{"me":  
{  
  "name": "Bernhard Hopfenmüller",  
  "occupation":  
    "Senior Consultant@ATIX AG",  
  "likes and works with":  
    "Automation,  
    Kafka,  
    Ansible,  
    Python,...",  
  "named a project once":  
    "Supreme Octo Potato",  
}  
}
```

Messaging-Systems

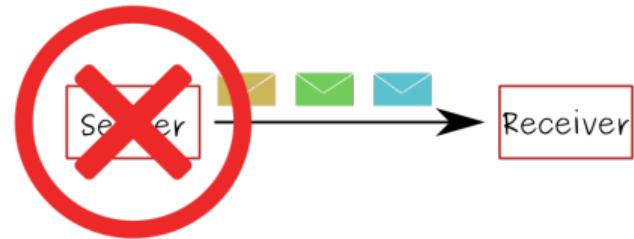
Why do we need a messaging system? [source 1]



Messaging Systems

Why do we need a messaging system? [source 1]

- ▶ Challenge 1: Sender not available
- ▶ Challenge 2: Sending too much (DoS)
- ▶ Challenge 3: Receiver crash upon processing



Messaging Systems

Why do we need a messaging system? [source 1]

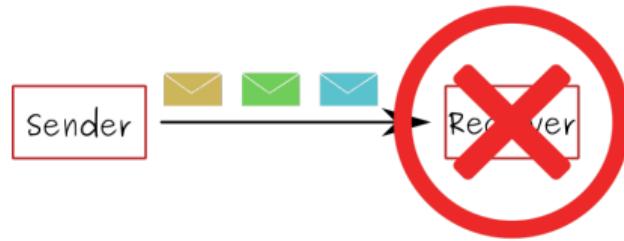
- ▶ Challenge 1: Sender not available
- ▶ Challenge 2: Sending too much (DoS)
- ▶ Challenge 3: Receiver crash upon processing



Messaging Systems

Why do we need a messaging system? [source 1]

- ▶ Challenge 1: Sender not available
- ▶ Challenge 2: Sending too much (DoS)
- ▶ Challenge 3: Receiver crash upon processing



Queues vs Topics

Waiting Queue vs Television

[source 1]

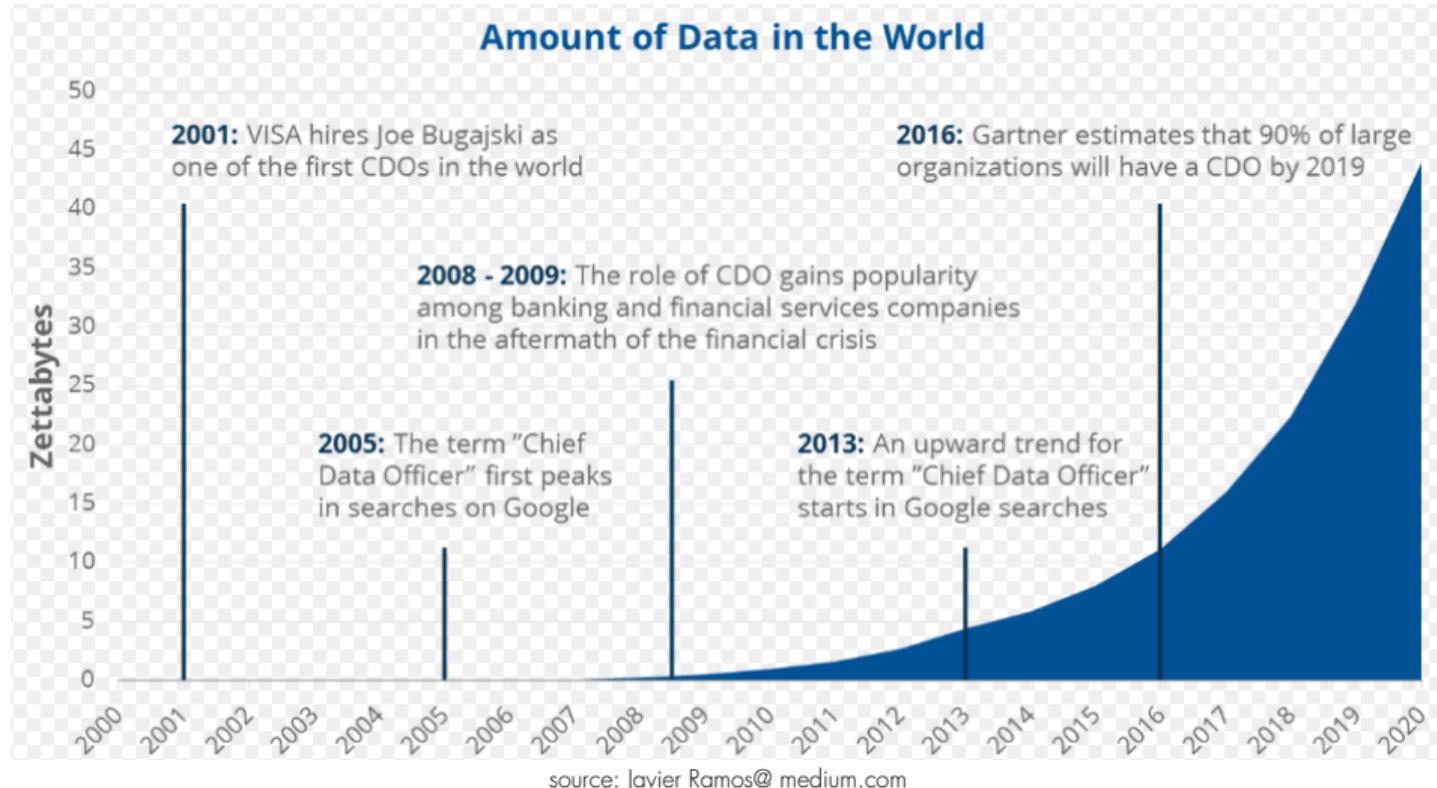
Wait until it's your turn



Choose what you want to receive



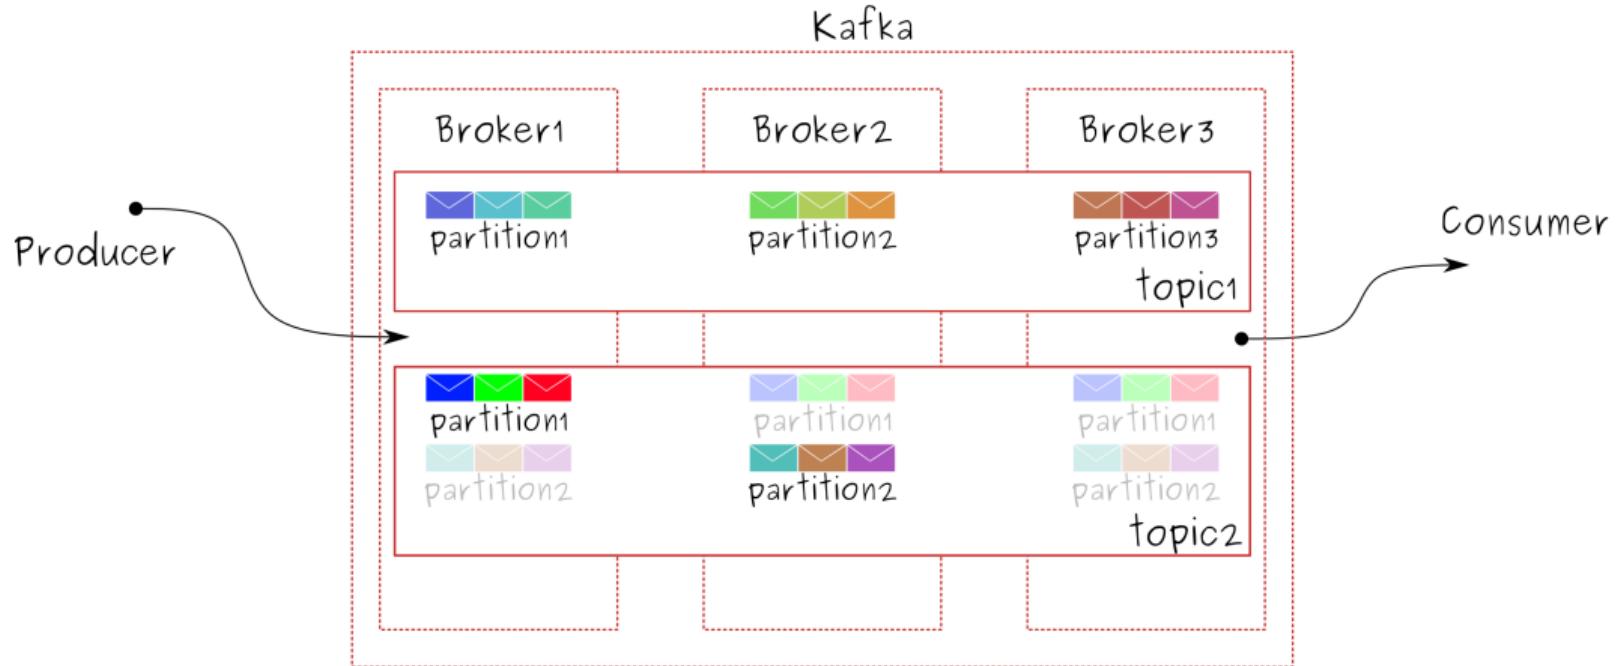
Why do we need “Television”



$$1\text{ZB} = 1000000000000000000000000000000000 \text{ bytes} = 10^{21} \text{ byte}$$

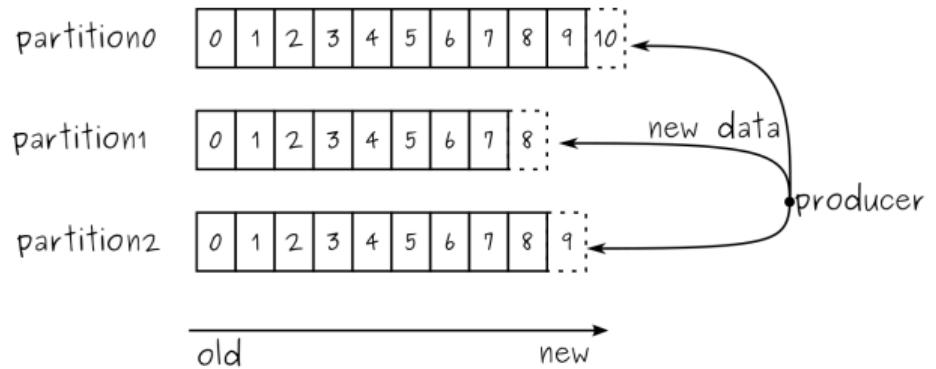
So let's watch some TV!

Kafka-Basic structure



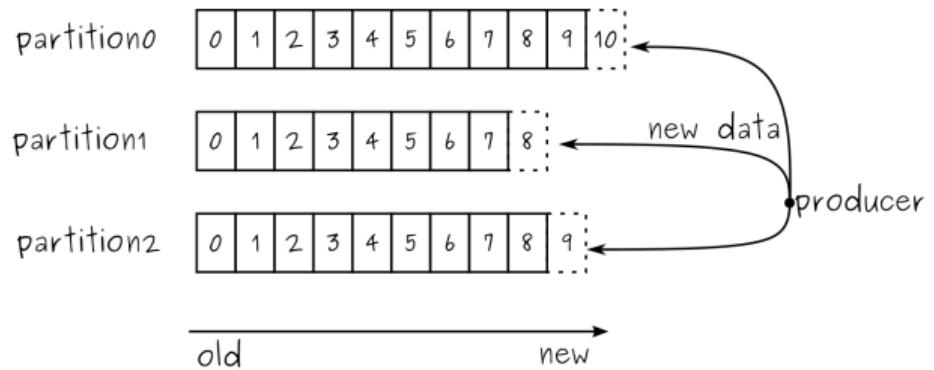
What is a topic in Kafka?

- ▶ core component of Kafka
- ▶ is filled by producer
- ▶ consists of one or more partitions



What is a topic in Kafka?

- ▶ producer can choose partition
- ▶ partition has running offset
- ▶ message is identified by offset



Data Storage in topics



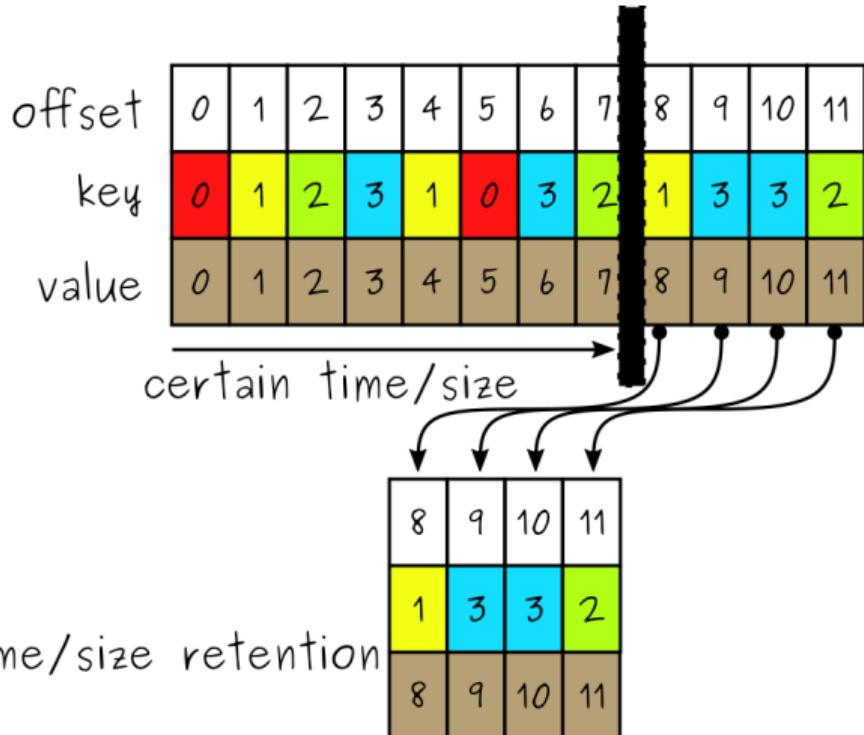
segment

0	4	8	12
---	---	---	----



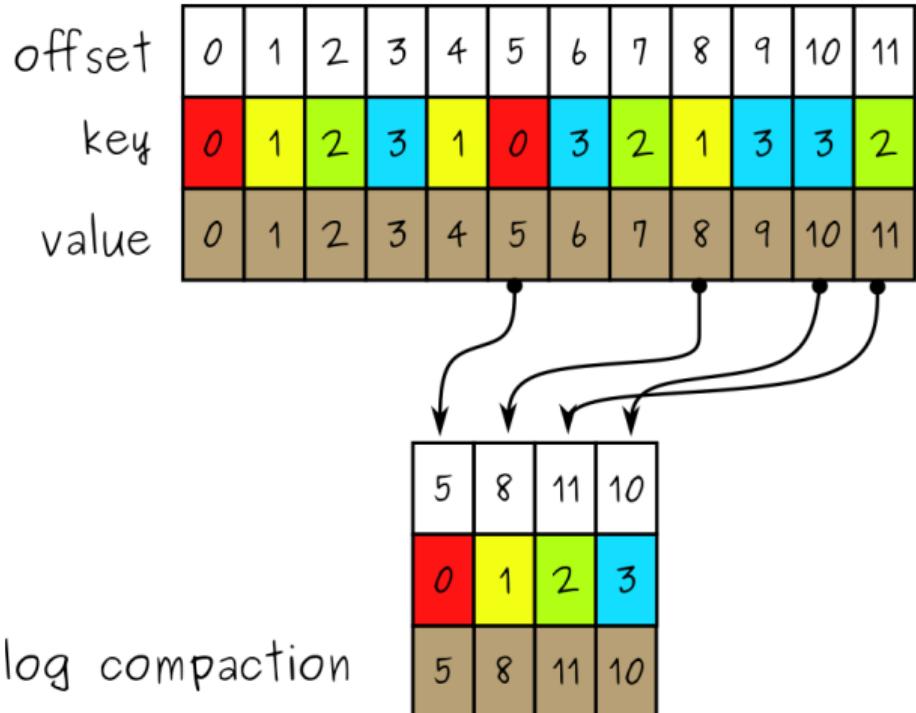
- ▶ messages are stored physically!
- ▶ key-value principle
- ▶ Clean-Up policies:

Cleaning up topics



- ▶ Clean-Up policies:
 - ▶ default: Retention-time
(delete old data after x days)
 - ▶ Retention-size
(delete old data if data memory $> x$)

Cleaning up topics

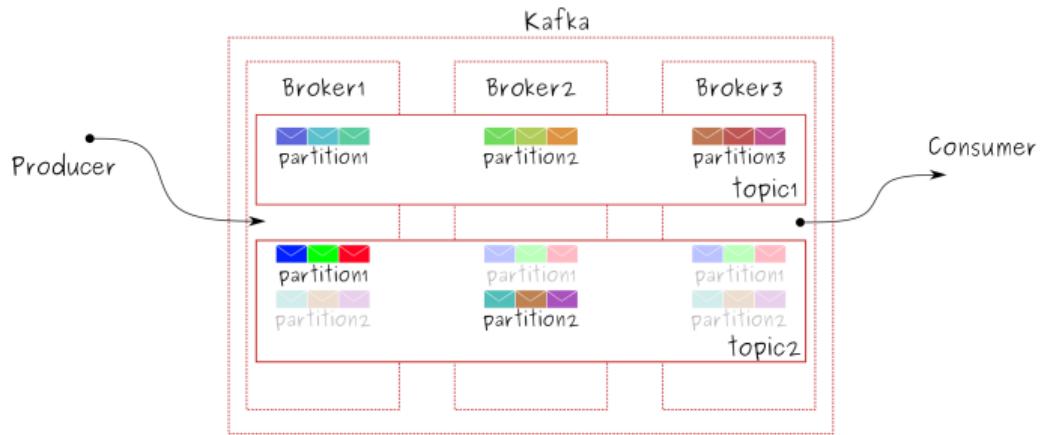


▶ Clean-Up policies:

- ▶ default: Retention-time
(delete old data after x days)
- ▶ Retention-size
(delete old data if data memory $> x$)
- ▶ Log-Compaction
(replace old value to key with new)

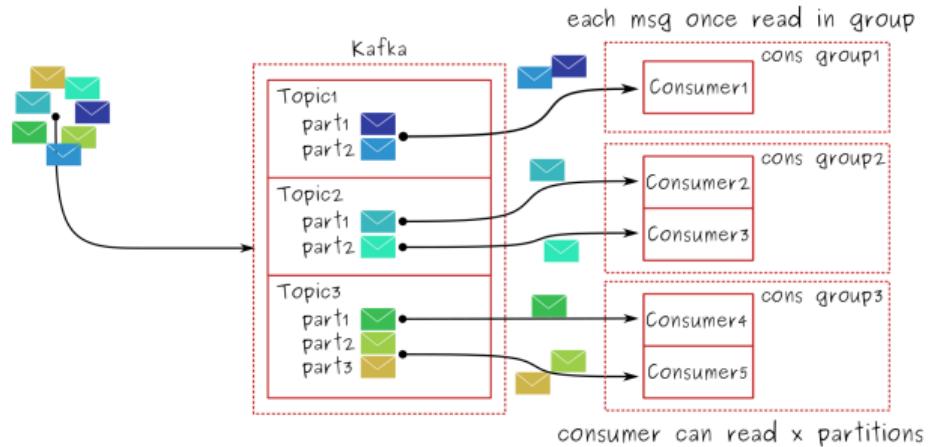
Topic consumption

- ▶ topics are pulled! (no DoS)
- ▶ any existing data can be pulled

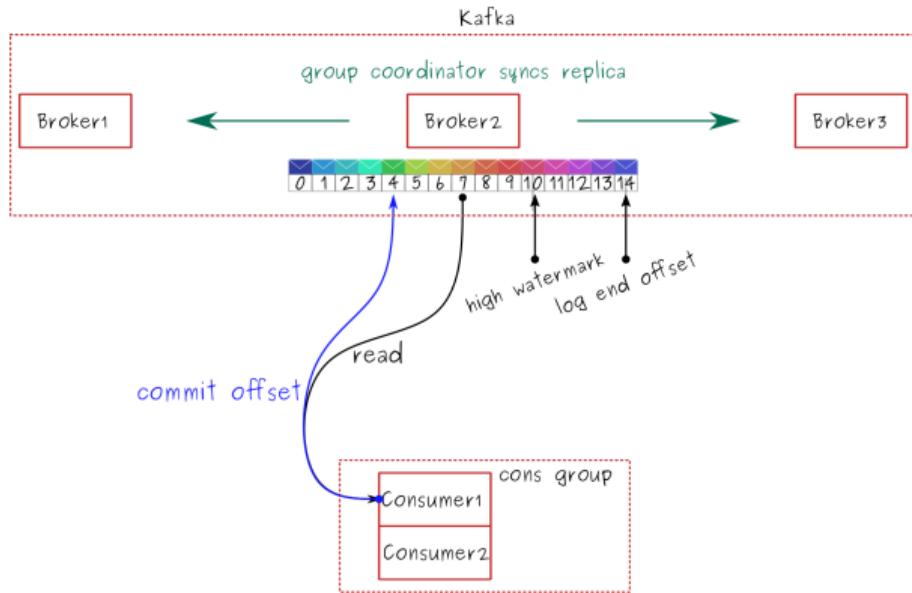


Consumer Groups

- ▶ parallelism allows high throughput
- ▶ never more consumers than partitions



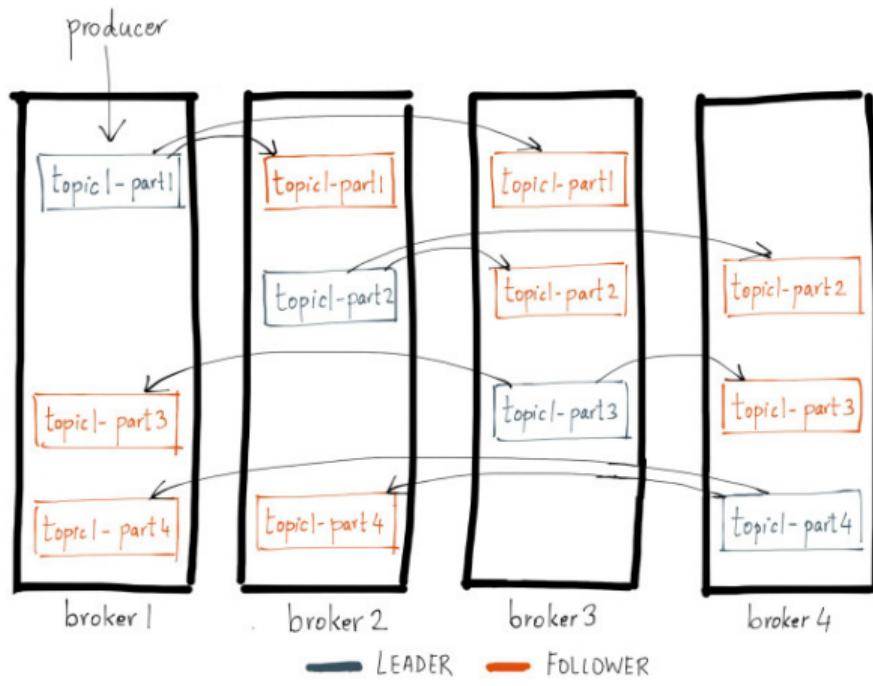
Wait but who knows what's read?



- ▶ Consumer commit their offset
- ▶ Upon failure re-processing possible

Replication

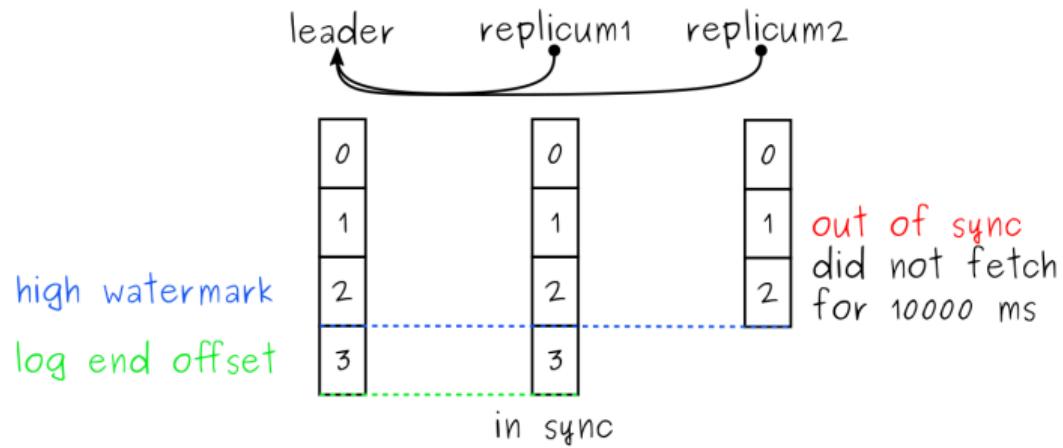
implemented on partition level



source: Confluent

In and Out of Sync Replica

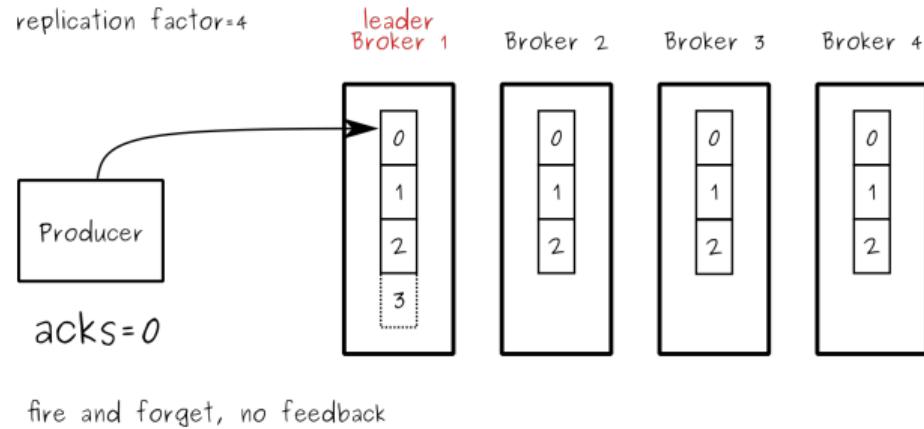
replica.lag.time.max.ms=10000



Did somebody hear my message?

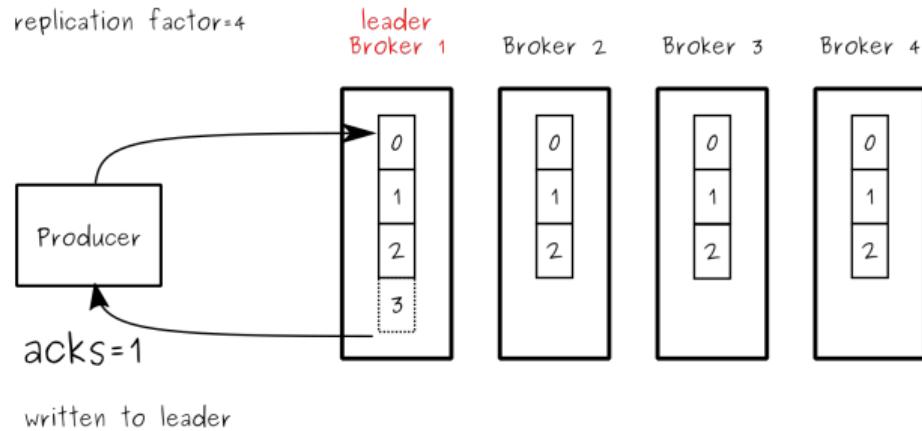
Producer decides if message was successfully sent

Delivery Guarantees



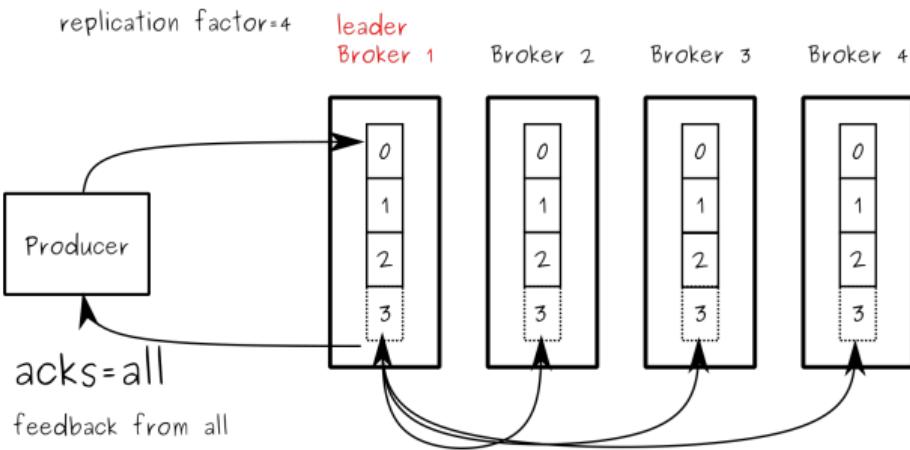
- ▶ as soon as sent
- ▶ as soon as received by first broker
- ▶ as soon as desired number of replica exist

Delivery Guarantees



- ▶ as soon as sent
- ▶ as soon as received by first broker
- ▶ as soon as desired number of replica exist

Delivery Guarantees



- ▶ as soon as sent
- ▶ as soon as received by first broker
- ▶ as soon as desired number of replica exist

Delivery Semantics

- ▶ At-most-once: Message is only delivered once or not at all
- ▶ At-least-once: Message is delivered once or more often
- ▶ Exactly-once: Message is delivered once - Kafka!

Delivery Semantics

- ▶ At-most-once: Message is only delivered once or not at all
- ▶ At-least-once: Message is delivered once or more often
- ▶ Exactly-once: Message is delivered once - Kafka!

Delivery Semantics

- ▶ At-most-once: Message is only delivered once or not at all
- ▶ At-least-once: Message is delivered once or more often
- ▶ Exactly-once: Message is delivered once - Kafka!

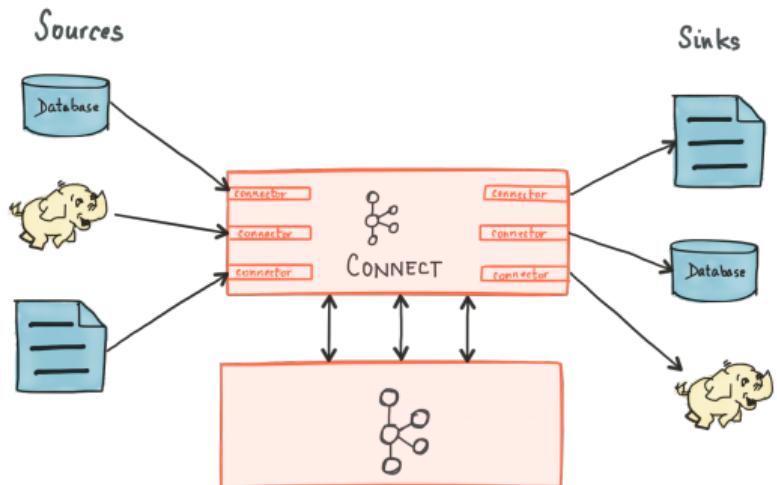
Green meadows?



Almost green meadows?



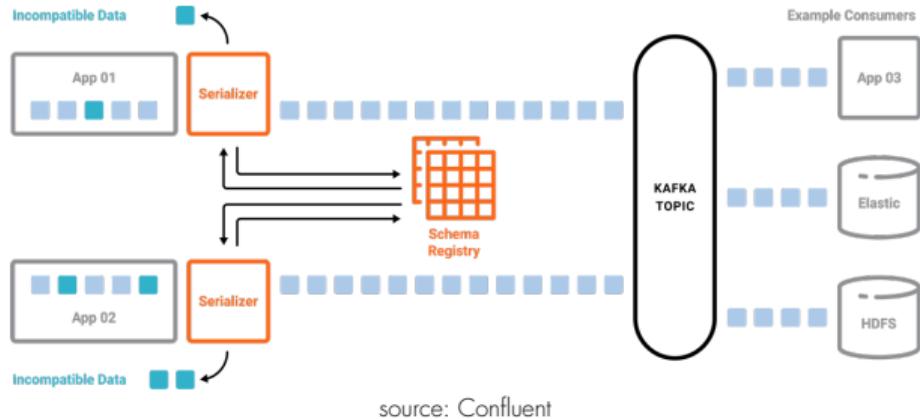
Talk to Kafka - Kafka Connect



- ▶ I/O for Kafka
- ▶ Connect with external/existing systems
- ▶ Open Source by Confluent

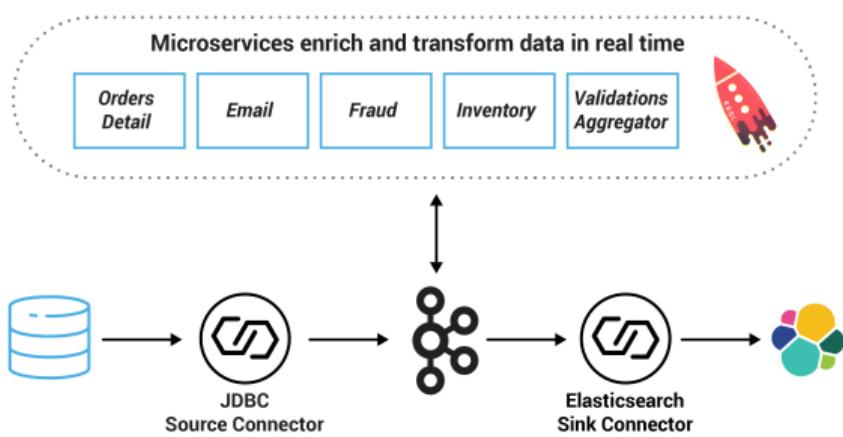
Talk to Kafka - Schema Registry

- ▶ define standards
- ▶ version and store them
- ▶ Open Source by Confluent



Going further: Kafka Streams

- ▶ reading data in real-time, process it, write it back
- ▶ scale applications
- ▶ Consumers is Producer
- ▶ different API



source: Confluent

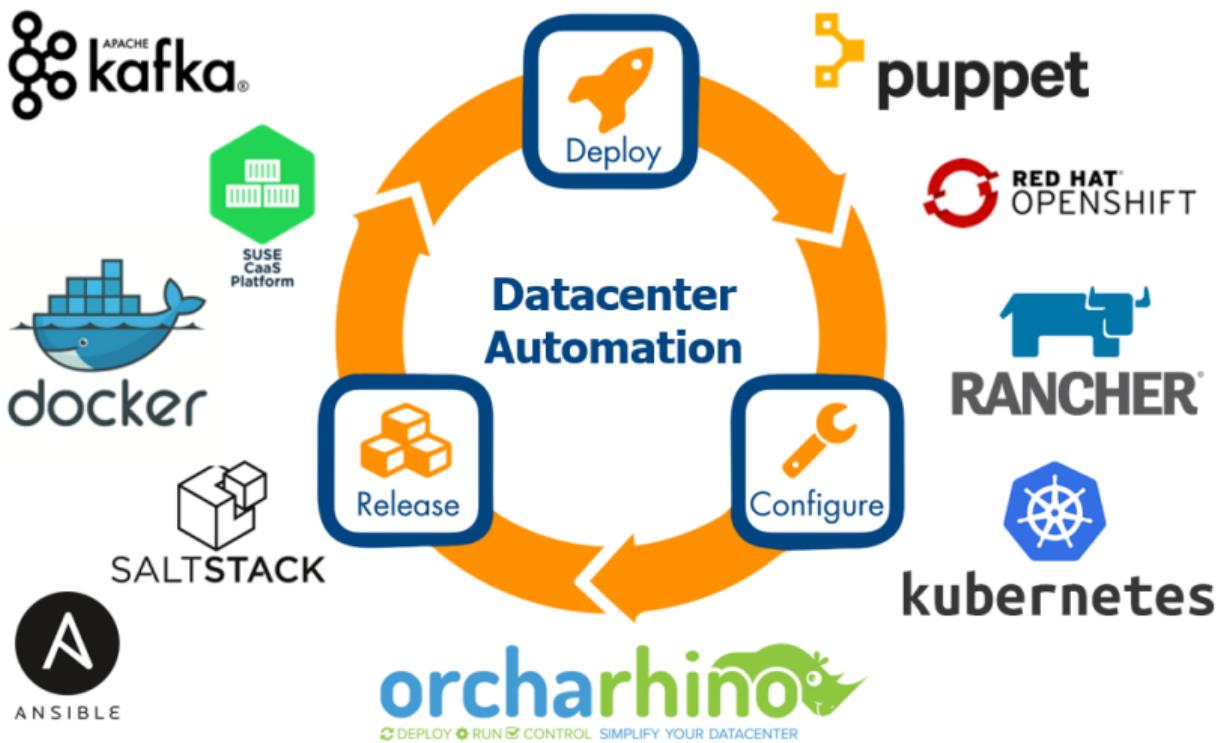
Use Cases

- ▶ Event Delivery/Streaming
- ▶ Website Activity Tracking
- ▶ Monitoring
- ▶ IoT-Messaging
- ▶ Storage

Read more

- ▶ ATIX Crew - <https://www.heise.de/select/ix/2019/4/1553935521978043>
- ▶ Lukas Berle [1] - <https://www.informatik-aktuell.de/betrieb/verfuegbarkeit/apache-kafka-eine-schlueselplattform-fuer-hochskalierbare-systeme.html>
- ▶ <https://kafka.apache.org/>
- ▶ <https://docs.confluent.io/current/>
- ▶ Google ;-)

Kafka is only part of the story!



curl -X GET localhost://whoami/contactme



```
{"me":  
 {  
 "xing/linkedin":  
 "Bernhard Hopfenmüller",  
 "email":  
 "hopfenmueller@atix.de",  
 "IRC/Github/Twitter":  
 "Fobhep",  
 }  
 }
```